

Notes on CUDA Programming

Aashay Kulkarni
December, 2025

Contents

1.	Introduction	2
1.1.	CUDA (Compute Unified Device Architecture)	2
2.	History of GPU Computing	3
2.1.	GPGPU	3
2.2.	GPU Computing	3
3.	Introduction to CUDA	4
3.1.	Data Parallelism	4
3.2.	CUDA Program Structure	4
3.2.1.	MATMUL Example	4
3.3.	Device Memory & Data Transfer	5
3.4.	Kernel Functions and Threading	7
4.	CUDA Threads	9
4.1.	CUDA Thread Organization	9

1. Introduction

Multicore CPUs enable you to maximize execution speed of sequential programs, while many-core GPUs allow greater execution throughput of parallel applications.

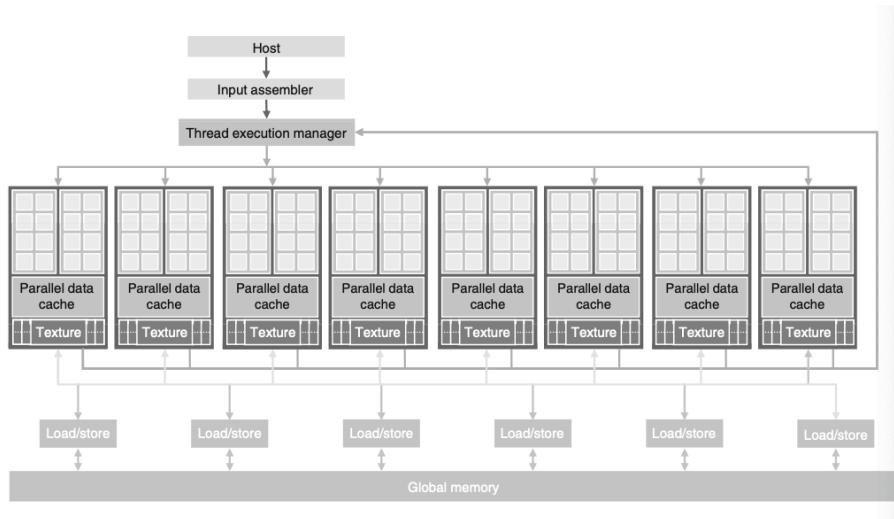
Today, there is a **large** performance gap between parallel and sequential execution.

Why?

- design of a CPU is optimized for sequential code performance
- memory bandwidth is another issue. GPUs move data much faster in and out of its DRAM.
- developers move computationally intensive parts of software to GPUs
- video games require massive number of floating point calculations per video frame, being executed in parallel and GPUs have been used for this purpose.

1.1. CUDA (Compute Unified Device Architecture)

- programming model created by NVIDIA to support joint CPU/GPU execution of an application.
- CUDA-capable GPU is organized into an array of highly threaded streaming multiprocessors (SMs). SMs combine to form a building block. SMs have a number of streaming processors (SPs) that share control logic and instruction cache.



- To experience speedup offered by parallelization, a large part of the application's execution time must be in the parallel portion.
- Certain applications have portions better suited to CPUs and hence a combined CPU/GPU parallel computing capability is required. This is precisely what CUDA promotes.
- Key steps in parallel computing:
 - identifying parts of application programs to be parallelized
 - isolating the data to be used by the parallelizing code by using API functions to allocate memory on the parallel computing device
 - using API functions to transfer data to the parallel computing device
 - developing kernel functions that will be executed by individual threads in the parallel part
 - launching kernel functions for execution by parallel parts
 - transferring data back to the host processor with API function calls

2. History of GPU Computing

2.1. GPGPU

- General Purpose Computing on GPUs.
- GPU processor array and frame buffer memory were designed to process graphics data and were too restrictive for general numerical applications.
- writes were extremely difficult -> could only be emitted as a pixel color value and configure the frame buffer operation to write.
- the handful of useful applications created with general computations on a GPU -> this field was called GPGPU.

2.2. GPU Computing

- NVIDIA developed Tesla GPU Architecture.
 - programming paradigm to think of GPU like a processor.
- programming approach involved explicit declaration of data-parallel aspects of their workload.
- no longer need to use graphics API to access parallel computing capabilities.

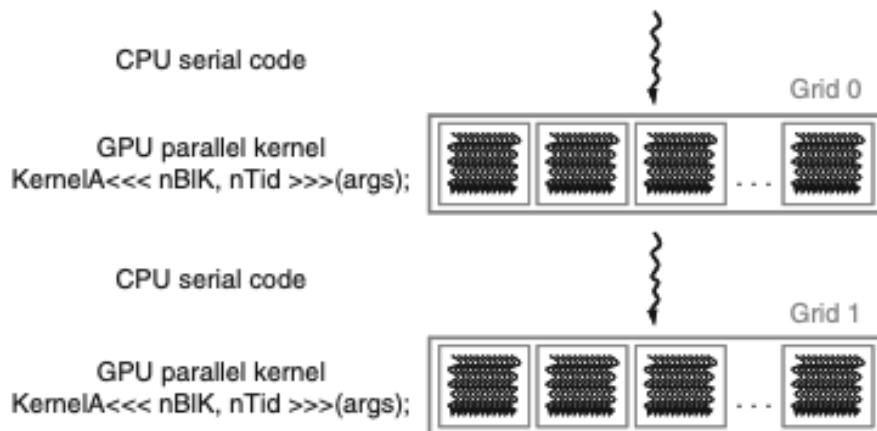
3. Introduction to CUDA

3.1. Data Parallelism

- computing system consists of host (CPU) & devices (massively parallel processors)
- CUDA devices accelerate execution of applications by harvesting a large amount of data parallelism.
- Matrix multiplication $P = M \times N$:
 - as every entry p_{ij} is independent of each other, a large amount of data parallelism can be performed.

3.2. CUDA Program Structure

- CUDA program comprises phases that are executed either by the host (CPU) or a device such as a GPU.
- CUDA program is a unified source code comprising both host & device code.
 - NVIDIA C compiler (nvcc): host code = ANSI C; device code = ANSI C extended with keywords for data-parallel functions, called kernels.
 - kernel functions generate a large number of threads to exploit data parallelism.
 - when no device is available, one can execute the kernel on a CPU using the CUDA SDK or MCUDA tool.
- CUDA threads are faster to generate and schedule than CPU threads due to efficient hardware support.



- CUDA program:
 - starts with CPU execution; when a kernel function is invoked, execution is moved to the device.
 - large numbers of threads are generated to take advantage of data parallelism, collectively called **grids**.
 - when all threads finish execution, the corresponding grid is terminated and execution moves back to the host.

3.2.1. MATMUL Example

```
int main (void) {
    // 1. Allocate and initialize matrices M, N, P
    //     I/O to read input matrices M and N

    // 2. M * N on the device
    //     MatrixMultiplication(M, N, P, Width);
```

```


    // 3. I/O to write output matrix P
    //     Free matrices M, N, P

    return 0;
}

```


3.3. Device Memory & Data Transfer

- Device code can:
 - R/W per-thread registers
 - R/W per-thread local memory
 - R/W per-block shared memory
 - R/W per-grid global memory
 - Read only per-grid constant memory
- Host code can
 - Transfer data to/from per-grid global and constant memories



- CUDA runtime system provides API functions to perform memory allocation and data transfer between host and devices.
- CUDA devices comprise global memory and constant memory; these are accessible from host code.
- Constant memory is read-only for device code.
- API functions `cudaMalloc()` and `cudaFree()` allocate and free global memory.
- API function `cudaMemcpy()` transfers data between host & device memory.

- `cudaMalloc()`
 - Allocates object in the device **global memory**
 - Two parameters
 - **Address of a pointer** to the allocated object
 - **Size of** of allocated object in terms of bytes
- `cudaFree()`
 - Frees object from device global memory
 - Pointer to freed object



For `cudaMalloc()`:

First parameter is a generic pointer (`void *`), Second parameter is the size in bytes

Example:

```

float *Md;
int size = Width * Width * sizeof(float);
cudaMalloc((void**)&Md, size);

// ...
cudaFree(Md);

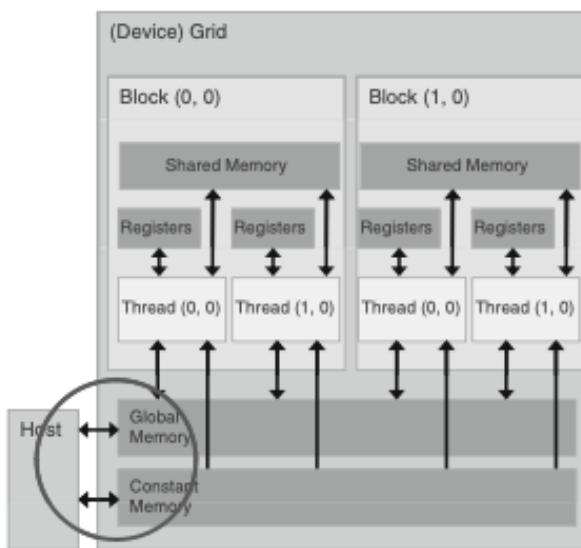
```

For cudaMemcpy():

First argument is a destination pointer, Second is a source pointer, Third is number of bytes and Fourth is direction (host→device, device→host, device→device)

Note: cudaMemcpy() cannot be used for memory transfer in multi-GPU systems.

- cudaMemcpy()
 - **Memory** data transfer
 - Requires four parameters
 - Pointer to destination
 - Pointer to source
 - Number of bytes copied
 - Type of transfer
 - Host to Host
 - Host to Device
 - Device to Host
 - Device to Device
 - Transfer is asynchronous



Example:

```

cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);

```

In the MatMul example, the main() function calls MatrixMultiplication(). MatrixMultiplication() allocates device memory, performs data transfers, and invokes the kernel that computes the result. This type of host-side function is called a stub function.

```

void MatrixMultiplication(float *M, float *N, float *P, int Width) {
    int size = Width * Width * sizeof(float);
    float *Md, Nd, Pd;

    // Allocate device memory for M, N, P
    cudaMalloc((void**)&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Pd, size);

    // Kernel invocation (not shown)
    // ...

    // Copy P & free device memory
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
}

```

```

    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}

```

3.4. Kernel Functions and Threading

In CUDA, a kernel function is executed by many threads in parallel. CUDA programming follows the single program, multiple data (SPMD) model.

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

Device functions cannot include recursion or indirect calls.

A function can be annotated for both host & device, generating two versions.

Variables like `threadIdx.x` and `threadIdx.y` give thread coordinates.


Thread notation: `thread_{threadIdx.x, threadIdx.y}`

Threads are organized hierarchically:

A grid contains one or more blocks

Each block has a unique 2D coordinate

- A thread block is a batch of threads that can cooperate with each other by:
 - Synchronizing their execution
 - For hazard-free shared memory accesses
 - Efficiently sharing data through a low-latency shared memory
- Two threads from two different blocks cannot cooperate



A block is a 3D array of threads (max 512 threads): indexed by `threadIdx.x, .y, .z`

When launching a kernel, host code chooses grid & block size.

Example:

```

// Setup execution configuration
dim3 dimBlock(Width, Width);
dim3 dimGrid(1, 1);

```

```
// Launch device computation threads!
c MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

4. CUDA Threads

4.1. CUDA Thread Organization

Threads in CUDA are organized in a two-level hierarchy using unique coordinates blockIdx and threadIdx. These are built-in, preinitialized variables, accessible in kernel functions.

In general, a grid is organized as a 2D array of blocks. Each block is organized into a 3D array of threads. The exact organization is determined by execution configuration.