

# Notes on CUDA Programming

Aashay Kulkarni  
December, 2025

## Contents

1.	Introduction .....	2
1.1.	CUDA (Compute Unified Device Architecture) .....	2
2.	History of GPU Computing .....	3
2.1.	GPGPU .....	3
2.2.	GPU Computing .....	3
3.	Introduction to CUDA .....	4
3.1.	Data Parallelism .....	4
3.2.	CUDA Program Structure .....	4
3.2.1.	MATMUL Example .....	4
3.3.	Device Memory & Data Transfer .....	5
3.4.	Kernel Functions and Threading .....	7
4.	CUDA Threads .....	8
4.1.	CUDA Thread Organization .....	8
4.2.	Using blockIdx AND threadIdx .....	9

# 1. Introduction

Multicore CPUs enable you to maximize execution speed of sequential programs, while many-core GPUs allow greater execution throughput of parallel applications.

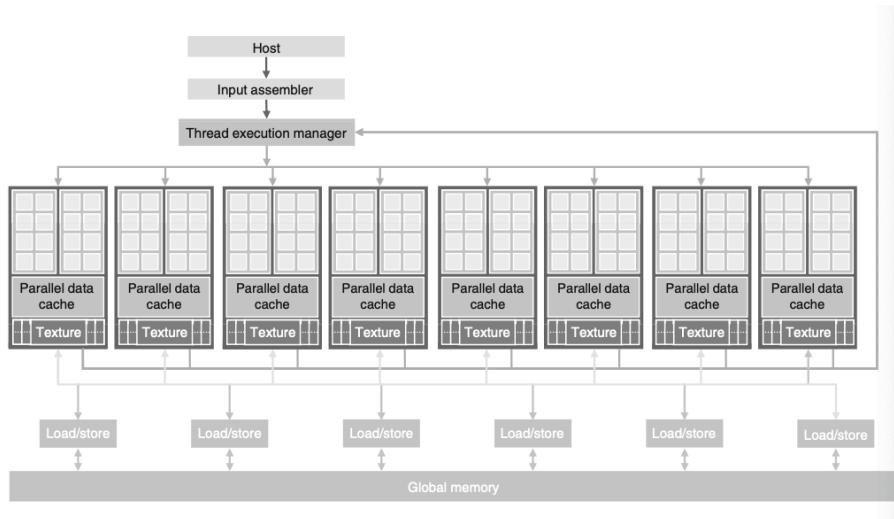
Today, there is a **large** performance gap between parallel and sequential execution.

**Why?**

- design of a CPU is optimized for sequential code performance
- memory bandwidth is another issue. GPUs move data much faster in and out of its DRAM.
- developers move computationally intensive parts of software to GPUs
- video games require massive number of floating point calculations per video frame, being executed in parallel and GPUs have been used for this purpose.

## 1.1. CUDA (Compute Unified Device Architecture)

- programming model created by NVIDIA to support joint CPU/GPU execution of an application.
- CUDA-capable GPU is organized into an array of highly threaded streaming multiprocessors (SMs). SMs combine to form a building block. SMs have a number of streaming processors (SPs) that share control logic and instruction cache.



- To experience speedup offered by parallelization, a large part of the application's execution time must be in the parallel portion.
- Certain applications have portions better suited to CPUs and hence a combined CPU/GPU parallel computing capability is required. This is precisely what CUDA promotes.
- Key steps in parallel computing:
  - identifying parts of application programs to be parallelized
  - isolating the data to be used by the parallelizing code by using API functions to allocate memory on the parallel computing device
  - using API functions to transfer data to the parallel computing device
  - developing kernel functions that will be executed by individual threads in the parallel part
  - launching kernel functions for execution by parallel parts
  - transferring data back to the host processor with API function calls

## **2. History of GPU Computing**

### **2.1. GPGPU**

- General Purpose Computing on GPUs.
- GPU processor array and frame buffer memory were designed to process graphics data and were too restrictive for general numerical applications.
- writes were extremely difficult -> could only be emitted as a pixel color value and configure the frame buffer operation to write.
- the handful of useful applications created with general computations on a GPU -> this field was called GPGPU.

### **2.2. GPU Computing**

- NVIDIA developed Tesla GPU Architecture.
  - programming paradigm to think of GPU like a processor.
- programming approach involved explicit declaration of data-parallel aspects of their workload.
- no longer need to use graphics API to access parallel computing capabilities.

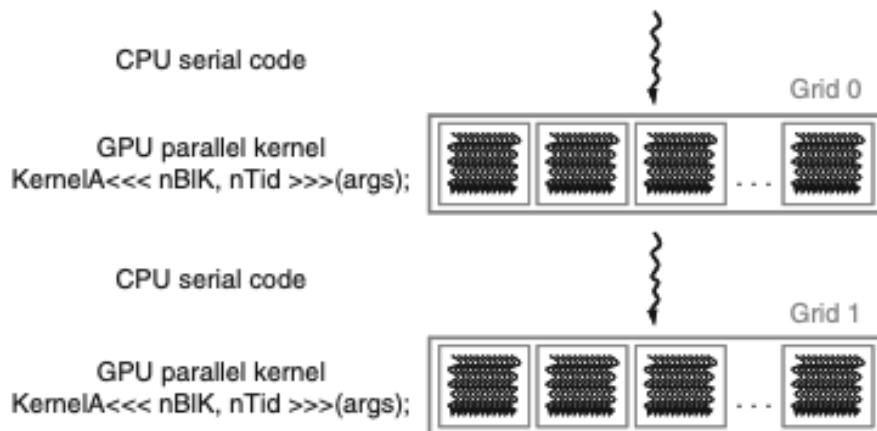
### 3. Introduction to CUDA

#### 3.1. Data Parallelism

- computing system consists of host (CPU) & devices (massively parallel processors)
- CUDA devices accelerate execution of applications by harvesting a large amount of data parallelism.
- Matrix multiplication  $P = M \times N$ :
  - as every entry  $p_{ij}$  is independent of each other, a large amount of data parallelism can be performed.

#### 3.2. CUDA Program Structure

- CUDA program comprises phases that are executed either by the host (CPU) or a device such as a GPU.
- CUDA program is a unified source code comprising both host & device code.
  - NVIDIA C compiler (nvcc): host code = ANSI C; device code = ANSI C extended with keywords for data-parallel functions, called kernels.
  - kernel functions generate a large number of threads to exploit data parallelism.
  - when no device is available, one can execute the kernel on a CPU using the CUDA SDK or MCUDA tool.
- CUDA threads are faster to generate and schedule than CPU threads due to efficient hardware support.



- CUDA program:
  - starts with CPU execution; when a kernel function is invoked, execution is moved to the device.
  - large numbers of threads are generated to take advantage of data parallelism, collectively called **grids**.
  - when all threads finish execution, the corresponding grid is terminated and execution moves back to the host.

##### 3.2.1. MATMUL Example

```
int main (void) {
    // 1. Allocate and initialize matrices M, N, P
    //     I/O to read input matrices M and N

    // 2. M * N on the device
    //     MatrixMultiplication(M, N, P, Width);
```

```

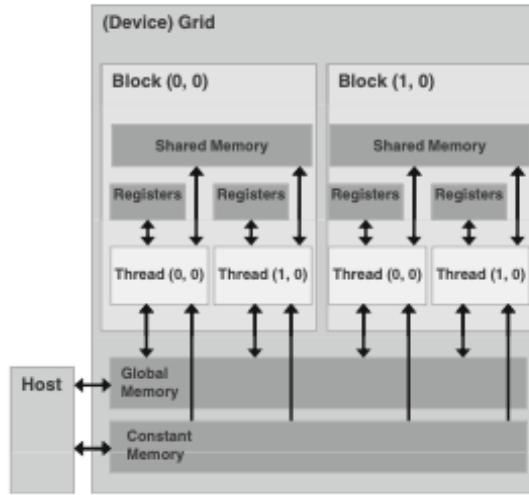
    // 3. I/O to write output matrix P
    //     Free matrices M, N, P

    return 0;
}

```

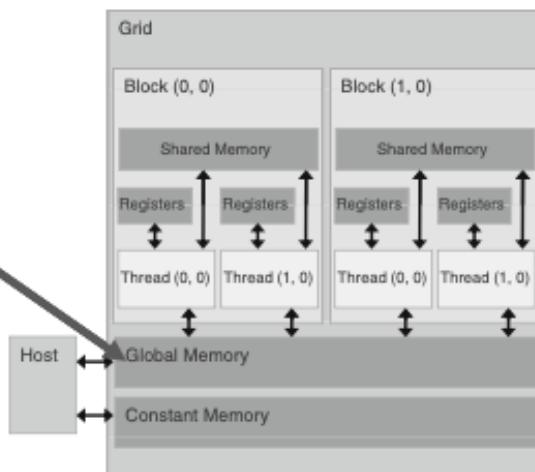
### 3.3. Device Memory & Data Transfer

- Device code can:
  - R/W per-thread registers
  - R/W per-thread local memory
  - R/W per-block shared memory
  - R/W per-grid global memory
  - Read only per-grid constant memory
- Host code can
  - Transfer data to/from per-grid global and constant memories



- CUDA runtime system provides API functions to perform memory allocation and data transfer between host and devices.
- CUDA devices comprise global memory and constant memory; these are accessible from host code.
- Constant memory is read-only for device code.
- API functions `cudaMalloc()` and `cudaFree()` allocate and free global memory.
- API function `cudaMemcpy()` transfers data between host & device memory.

- `cudaMalloc()`
  - Allocates object in the device **global memory**
  - Two parameters
    - **Address of a pointer** to the allocated object
    - **Size of** of allocated object in terms of bytes
- `cudaFree()`
  - Frees object from device global memory
    - Pointer to freed object



#### For `cudaMalloc()`:

First parameter is a generic pointer (`void *`), Second parameter is the size in bytes

Example:

```

float *Md;
int size = Width * Width * sizeof(float);
cudaMalloc((void**)&Md, size);

// ...
cudaFree(Md);

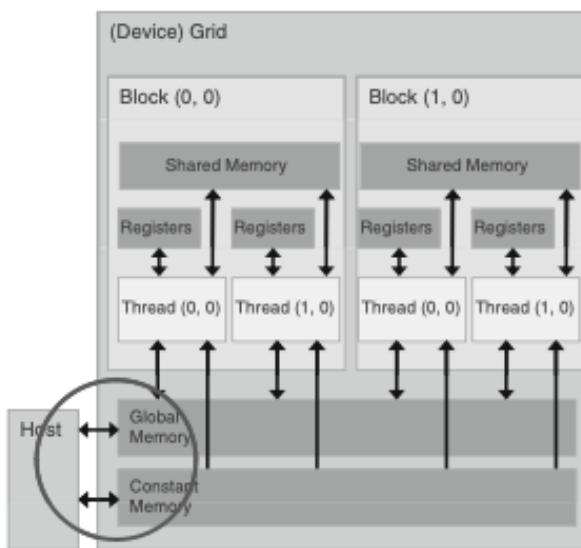
```

### For cudaMemcpy():

First argument is a destination pointer, Second is a source pointer, Third is number of bytes and Fourth is direction (host→device, device→host, device→device)

Note: cudaMemcpy() cannot be used for memory transfer in multi-GPU systems.

- cudaMemcpy()
  - **Memory** data transfer
  - Requires four parameters
    - Pointer to destination
    - Pointer to source
    - Number of bytes copied
  - Type of transfer
    - Host to Host
    - Host to Device
    - Device to Host
    - Device to Device
  - Transfer is asynchronous



Example:

```

cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);

```

In the MatMul example, the main() function calls MatrixMultiplication(). MatrixMultiplication() allocates device memory, performs data transfers, and invokes the kernel that computes the result. This type of host-side function is called a stub function.

```

void MatrixMultiplication(float *M, float *N, float *P, int Width) {
    int size = Width * Width * sizeof(float);
    float *Md, Nd, Pd;

    // Allocate device memory for M, N, P
    cudaMalloc((void**)&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Pd, size);

    // Kernel invocation (not shown)
    // ...

    // Copy P & free device memory
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
}

```

```

    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}

```

### 3.4. Kernel Functions and Threading

In CUDA, a kernel function is executed by many threads in parallel. CUDA programming follows the single program, multiple data (SPMD) model.

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

Device functions cannot include recursion or indirect calls.

A function can be annotated for both host & device, generating two versions.

Variables like `threadIdx.x` and `threadIdx.y` give thread coordinates.

Thread notation: `thread{threadIdx.x, threadIdx.y}`

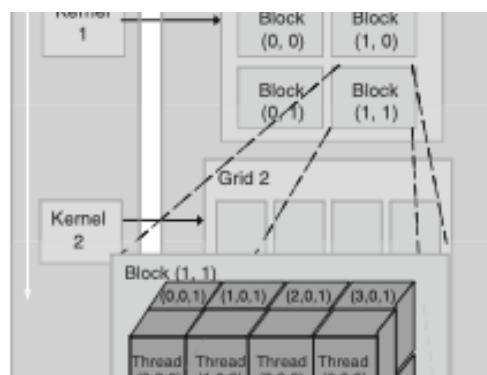
Threads are organized hierarchically:

A grid contains one or more blocks

Each block has a unique 2D coordinate

other by:

- Synchronizing their execution
  - For hazard-free shared memory accesses
- Efficiently sharing data through a low-latency shared memory
  - Two threads from two different blocks cannot



A block is a 3D array of threads (max 512 threads): indexed by `threadIdx.x`, `.y`, `.z`

When launching a kernel, host code chooses grid & block size.

Example:

```

// Setup execution configuration
dim3 dimBlock(Width, Width);
dim3 dimGrid(1, 1);

// Launch device computation threads!
c MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);

```

## 4. CUDA Threads

### 4.1. CUDA Thread Organization

Threads in CUDA are organized in a two-level hierarchy using unique coordinates blockIdx and threadIdx. These are built-in, preinitialized variables, accessible in kernel functions.

In general, a CUDA grid is organized as a 2D array of blocks. Each block is organized into a 3D array of threads. The exact organization is determined by execution configuration.

Kernels are launched with a given execution parameter of a dim3 type which is a C struct with three unsigned int fields: x,y,z.

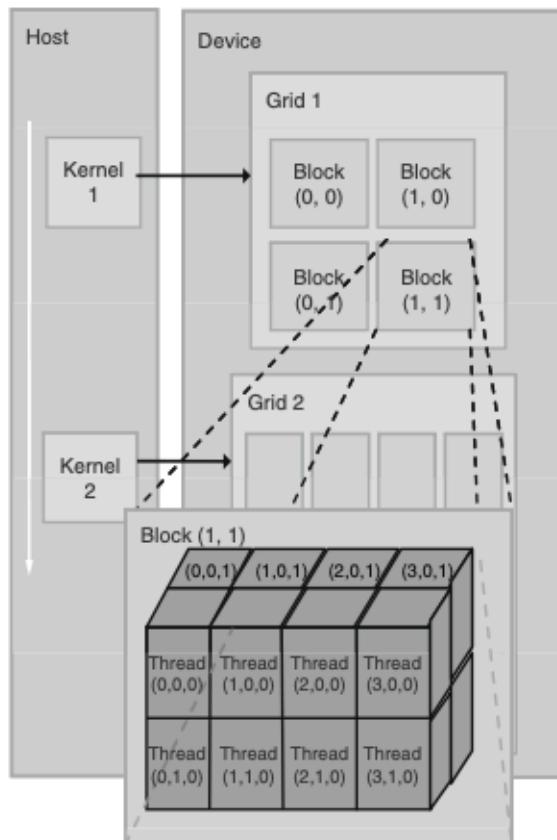
Example:

```
dim3 dimGrid(128, 1, 1)
dim3 dimBlock(32, 1, 1)
KernelFunction<<<dimGrid, dimBlock>>>(...);
```

If a grid / block only has one dimension, scalar values can be used.

```
KernelFunction<<<128, 32>>>(...);
```

Values of gridDims go from 1 to 65535. Values of blockIdx go from 0 to gridDim - 1. A *threadIdx* has an x,y,z coordinate. Each block is limited to 1024 threads (it was 512 earlier).



**FIGURE 4.2**

A multidimensional example of CUDA grid organization.

## **4.2. Using blockIdx AND threadIdx**