

# Programming Massively Parallel Processors

## CH1: Introduction


- multicore CPUs: maximize execution speed of sequential programs
- many-core GPUs: execution throughput of parallel applications
- large performance gap between parallel and sequential execution

Why?

- developers move computationally intensive parts of software to GPUs
- design of a CPU is optimized for sequential code performance
- memory bandwidth is another issue. GPUs move data much faster in and out of its DRAM.
- video games require massive number of floating point calculations per video frame, being executed in parallel and GPUs have been used for this purpose.

CUDA (Compute Unified Device Architecture)

- programming model created by NVIDIA to support joint CPU/GPU execution of an application.
- CUDA-capable GPU is organized into an array of highly threaded streaming multiprocessors (SMs). SMs combine to form a building block. SMs have a number of streaming processors (SPs) that share control logic and instruction cache.



- To experience speedup offered by parallelization, a large part of the application's execution time must be in the parallel portion.

- Certain applications have portions better suited to CPUs and hence a combined CPU/GPU parallel computing capability is required. This is precisely what CUDA promotes.
  - Key steps in parallel computing:
    - identifying parts of application programs to be parallelized
    - isolating the data to be used by the parallelizing code by using API functions to allocate memory on the parallel computing device
    - using API functions to transfer data to the parallel computing device
    - developing kernel functions that will be executed by individual threads in the parallel part
    - launching kernel functions for execution by parallel parts
    - transferring data back to the host processor with API function calls
- 

## CH2: History of GPU Computing

### GPGPU

General Purpose Computing on GPUs.

- GPU processor array and frame buffer memory were designed to process graphics data and were too restrictive for general numerical applications.
- writes were extremely difficult → could only be emitted as a pixel color value and configure the frame buffer operation to write.
- the handful of useful applications created with general computations on a GPU → this field was called GPGPU.

### GPU Computing

- NVIDIA developed Tesla GPU Architecture.
  - programming paradigm to think of GPU like a processor.
  - programming approach involved explicit declaration of data-parallel aspects of their workload.
  - no longer need to use graphics API to access parallel computing capabilities.
- 

## CH3: Introduction to CUDA


### Data Parallelism

- computing system consists of host (CPU) & devices (massively parallel processors)

- CUDA devices accelerate execution of applications by harvesting a large amount of data parallelism.
- Matrix multiplication  $P = M \times N$ :
  - as every entry  $p_{ij}$  is independent of each other, a large amount of data parallelism can be performed.

### CUDA Program Structure

- CUDA program comprises phases that are executed either by the host (CPU) or a device such as a GPU.
- CUDA program is a unified source code comprising both host & device code.
  - NVIDIA C compiler (nvcc): host code = ANSI C; device code = ANSI C extended with keywords for data-parallel functions, called kernels.
  - kernel functions generate a large number of threads to exploit data parallelism.
  - when no device is available, one can execute the kernel on a CPU using the CUDA SDK or MCUDA tool.
- CUDA threads are faster to generate and schedule than CPU threads due to efficient hardware support.




- CUDA program:
  - starts with CPU execution; when a kernel function is invoked, execution is moved to the device.
  - large numbers of threads are generated to take advantage of data parallelism, collectively called grids.
  - when all threads finish execution, the corresponding grid is terminated and execution moves back to the host.

## MATMUL Example

```
int main (void) {  
    // 1. Allocate and initialize matrices M, N, P  
    //     I/O to read input matrices M and N  
  
    // 2. M * N on the device  
    //     MatrixMultiplication(M, N, P, Width);  
  
    // 3. I/O to write output matrix P  
    //     Free matrices M, N, P  
  
    return 0;  
}
```

## Device Memory & Data Transfer

- Device code can:
  - R/W per-thread registers
  - R/W per-thread local memory
  - R/W per-block shared memory
  - R/W per-grid global memory
  - Read only per-grid constant memory
- Host code can
  - Transfer data to/from per-grid global and constant memories



CUDA runtime system provides API functions to perform memory allocation and data transfer between host and devices.


CUDA devices comprise global memory and constant memory; these are accessible from host code.

Constant memory is read-only for device code.

API functions `cudaMalloc()` and `cudaFree()` allocate and free global memory.

API function `cudaMemcpy()` transfers data between host & device memory.

- `cudaMalloc()`
  - Allocates object in the device global memory
  - Two parameters
    - **Address of a pointer** to the allocated object
    - **Size of** of allocated object in terms of bytes
- `cudaFree()`
  - Frees object from device global memory
    - Pointer to freed object



For `cudaMalloc()`:

First parameter is a generic pointer (`void`)

Second parameter is the size in bytes

Example:

```

float *Md;
int size = Width * Width * sizeof(float);
cudaMalloc((void**)&Md, size);

// ...
cudaFree(Md);
  
```

For `cudaMemcpy()`:

First argument = destination pointer


Second = source pointer

Third = number of bytes

Fourth = direction (host→device, device→host, device→device)

Note: `cudaMemcpy()` cannot be used for memory transfer in multi-GPU systems.

- \* `cudaMemcpy()`
  - **Memory** data transfer
  - Requires four parameters
    - Pointer to destination
    - Pointer to source
    - Number of bytes copied
  - Type of transfer
    - Host to Host
    - Host to Device
    - Device to Host
    - Device to Device
  - Transfer is asynchronous



### Example:

```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
```

In the MatMul example, the `main()` function calls `MatrixMultiplication()`. `MatrixMultiplication()` allocates device memory, performs data transfers, and invokes the kernel that computes the result. This type of host-side function is called a stub function.

```
void MatrixMultiplication(float *M, float *N, float *P, int Width) {
    int size = Width * Width * sizeof(float);
    float *Md, Nd, Pd;

    // Allocate device memory for M, N, P
    cudaMalloc((void**)&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Pd, size);

    // Kernel invocation (not shown)
    // ...

    // Copy P & free device memory
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

## Kernel Functions and Threading

In CUDA, a kernel function is executed by many threads in parallel. CUDA programming follows the single program, multiple data (SPMD) model.

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

Device functions cannot include recursion or indirect calls.

A function can be annotated for both host & device, generating two versions.

Variables like `threadIdx.x` and `threadIdx.y` give thread coordinates.


Thread notation: `thread{threadIdx.x, threadIdx.y}`

Threads are organized hierarchically:

A grid contains one or more blocks

Each block has a unique 2D coordinate

- A thread block is a batch of threads that can cooperate with each other by:
  - Synchronizing their execution
    - For hazard-free shared memory accesses
  - Efficiently sharing data through a low-latency shared memory
- Two threads from two different blocks cannot cooperate



A block is a 3D array of threads (max 512 threads): indexed by `threadIdx.x, .y, .z`

When launching a kernel, host code chooses grid & block size.

Example:

```
dim3 dimBlock(Width, Width); dim3 dimGrid(1, 1);
```

```
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

## CH4: CUDA Threads